

The ANARA Production Benchmark 2026

Evaluating Frontier AI Coding Models Against a Live Production AI System

Version 1.0 - Public Release

Generated from evidence-labeled benchmark reports and structured datasets. Results, scores, rankings, and conclusions are unchanged from the completed benchmark package.

Executive Summary

Most AI coding benchmarks ask models to solve isolated programming tasks. That is useful, but it is not the same thing as production engineering. Real release decisions depend on whether a model can connect source code, deployment state, runtime configuration, provider behavior, security controls, and live user experience without confusing local fixes for deployed behavior.

The ANARA Production Benchmark 2026 exists to test that harder question. It evaluates frontier AI coding models against the ANARA production voice platform: a live Cloud Run system with Vapi integration, a multi-provider LLM chain, phone identity mapping, weather, Gmail, Calendar, provider diagnostics, and live phone-call acceptance testing.

ANARA was chosen because it resembles the systems engineering teams actually ship: stateful, provider-dependent, security-sensitive, and exposed to real users. A synthetic coding task can show whether a model writes code. ANARA shows whether a model can reason about production readiness.

The major finding is clear and unchanged: production was not ready for another Wingman call. Five completed forensic audits independently converged on critical blockers, and the Wingman live validation confirmed a failed phone experience. Claude Fable 5 is marked incomplete and excluded from comparative rankings because its benchmark session did not complete.

Table 1. Executive summary of unchanged benchmark findings.

Finding	What The Evidence Shows	Verification
Public LLM endpoint exposure	`/v1/chat/completions` accepted unauthenticated production requests.	Runtime Verified
Caller / owner spoofing	Caller metadata supplied in the request body could reach owner-aware behavior.	Runtime Verified
Deployment drift	Reported VAPI fixes were not fully reflected in the serving revision.	Source-Code Verified
Provider reality	Hugging Face and GitHub were live; OpenAI invalid; Azure/Gemini absent.	Runtime Verified
Wingman failure	Live phone validation produced repeated response-path stalls and hangup-loop failure.	Manual Verification

Key Findings At A Glance

This page is designed for two-minute review. It separates model status, scoring, and consensus findings without changing rankings or conclusions.

Table 2. Models tested, final status, score availability, and ranking treatment.

Model	Status	Readiness	Benchmark score	Ranking treatment
GPT-5.5	extracted sanitized	5/10	68.7	Included where complete
GPT-5.3-Codex / GPT-5.3-Code	extracted sanitized	4/10	Not reported	Included where complete
Claude Opus 4.8 (fast mode) (Preview)	extracted sanitized	5/10	Not reported	Included where complete
DeepSeek V4 Pro (thinking)	extracted sanitized	5/10	Not reported	Included where complete
Gemini 3.1 Pro (Preview)	extracted sanitized	4/10	35.9	Included where complete
Claude Fable 5	incomplete session instability	Not assigned	Not reported	exclude_from_final_comparative_rankings

Model	Status	Readiness	Benchmark score	Ranking treatment
Wingman Validation	extracted sanitized	Not assigned	2/10	Included where complete

Table 3. Consensus findings at a glance.

Consensus Finding	Status	Evidence Type
Authentication bypass	Confirmed by completed model audits	Runtime Verified
Deployment drift	Confirmed by completed model audits	Source-Code Verified
Owner spoofing	Confirmed by completed model audits	Runtime Verified
OpenAI credential issue	Confirmed by completed model audits	Runtime Verified
Weather degradation / subsystem behavior	Reported as working or gracefully degraded	Source-Code / Runtime Verified
Provider diagnostics	Reported present	Source-Code / Runtime Verified
Wingman validation failure	Confirmed by live call validation	Manual Verification

Verification Legend

[Runtime Verified] observed through live endpoint, provider probe, Cloud Run state, logs, or phone-call evidence.

[Source-Code Verified] observed through source trace, git diff, or repository inspection reported by a benchmark run.

[Manual Verification] observed directly in supplied human or Wingman validation evidence.

[AI Interpretation] model synthesis or interpretation derived from the evidence and kept distinct from verified facts.

Why Production Benchmarks Matter

Synthetic coding benchmarks are controlled and repeatable, but they tend to compress engineering into a single prompt-response loop. Production systems do not behave that way. A real deployment can fail because a route is mounted outside middleware, an environment variable is stale, a provider key is invalid, a dirty working tree diverges from the serving image, or a live call enters a repetition loop.

This benchmark therefore measures production reasoning rather than raw coding ability. It rewards evidence discipline: models must state what was runtime verified, what was source-code verified, what was manually verified, and what remains interpretation.

Methodology

Every forensic model received the same objective, the same production target, and the same class of task: audit ANARA production readiness without fabricating results. The target was the ANARA production deployment and its Vapi/LLM/phone path.

- Identical prompt class: each model was asked to perform a production-readiness investigation.
- Identical production target: each completed model evaluated the same ANARA deployment context.
- Identical objective: determine whether ANARA was ready for another Wingman production call.
- Evidence labeling: every extracted finding records model, evidence, confidence, category, verification status, and report location.

- No-fabrication rule: incomplete runs are marked incomplete rather than scored.

Why Fable 5 Is Excluded

Claude Fable 5 did not complete the benchmark. The run showed promising reasoning before failure, but the investigation terminated before independent production findings could be verified. Ranking it last would invent a result.

Version 1.0 therefore lists Fable as Incomplete (Session Instability) and excludes it from comparative rankings.

Why Wingman Is Acceptance Testing

Wingman is not treated as a peer model forensic audit. It is live production acceptance testing: a natural phone conversation that tests whether the deployed system works for a user. Its 2/10 score validates production readiness from the outside-in.

Results

Table 4. Unchanged readiness and validation scores.

Model	Readiness / Score	Verdict
Gemini 3.1 Pro (Preview)	4/10	STILL BLOCKED
DeepSeek V4 Pro (thinking)	5/10	STILL BLOCKED
Claude Opus 4.8 (fast mode) (Preview)	5/10	STILL BLOCKED
GPT-5.5	5/10	STILL BLOCKED
GPT-5.3-Codex / GPT-5.3-Code	4/10	STILL BLOCKED
Wingman Validation	2/10	NOT READY FOR PRODUCTION
Claude Fable 5	Not scored	Incomplete

Consensus Findings

[Runtime Verified] ``v1/chat/completions`` is unauthenticated and reachable without a secret header. Evidence: All five reports describe direct runtime probes returning HTTP 200 from ``v1/chat/completions`` without authentication.

[Runtime Verified] Caller or owner identity can be spoofed through request body caller metadata. Evidence: All five reports identify request-body caller metadata as trusted client-supplied input that can spoof owner context. Personal identifiers redacted.

[Source-Code Verified] Deployment drift exists: important VAPI fixes are uncommitted and not deployed. Evidence: All five reports identify dirty working-tree deployment drift with important VAPI hardening or owner-mapping changes not fully verifiable in the serving image.

[Runtime Verified] OpenAI is unavailable due to invalid API key, while Hugging Face and GitHub Models are live. Evidence: All five reports cite OpenAI 401 and successful Hugging Face/GitHub provider probes.

[AI Interpretation] Final production verdict is STILL BLOCKED. Evidence: All five reports end with a STILL BLOCKED verdict.

[Manual Verification] Independent live validation agrees production should not proceed. Evidence: Five model audits conclude STILL BLOCKED; Wingman natural conversation validation concludes NOT READY FOR PRODUCTION with a 2/10 score.

Security Findings

Table 5. Highest-priority security findings.

Severity	Model	Category	Evidence	Verification
CRITICAL	Gemini 3.1 Pro (Preview)	unauthenticated_endpoint	`POST /v1/chat/completions` and `POST /api/vapi/llm` are reported publicly exposed with no authentication.	runtime_verified
CRITICAL	Gemini 3.1 Pro (Preview)	identity_spoofing	The report says caller identity is spoofable via `body.call.customer.number`, enabling unauthorized owner context mapping. Personal identifier redacted.	runtime_verified
CRITICAL	DeepSeek V4 Pro (thinking)	unauthenticated_endpoint	`/v1/chat/completions` is reported to have no authentication and caller ID is spoofable.	runtime_verified
CRITICAL	DeepSeek V4 Pro (thinking)	unauthenticated_endpoint	`/api/vapi/llm` is reported to share the same vulnerability and to be in `PUBLIC_API_ROUTES`.	source_code_verified
CRITICAL	Claude Opus 4.8 (fast mode) (Preview)	unauthenticated_endpoint	`/v1/chat/completions` and `/api/vapi/llm` are reported to accept unauthenticated requests.	runtime_verified
CRITICAL	Claude Opus 4.8 (fast mode) (Preview)	identity_spoofing	The report says caller identity and owner context are spoofable through request-body phone metadata, creating a private-context injection/exfiltration vector. Personal identifier redacted.	runtime_verified
CRITICAL	GPT-5.5	unauthenticated_endpoint	The report says `POST /v1/chat/completions` returned HTTP 200 without `x-vapi-secret`.	runtime_verified
CRITICAL	GPT-5.5	unauthenticated_endpoint	The report says `/api/vapi/llm` shares the same handler and lacks `verifyVapiRequest`.	source_code_verified
CRITICAL	GPT-5.5	identity_spoofing	The report says body-provided `call.customer.number` is trusted and spoofing the owner phone number reaches owner-aware behavior. Personal identifier redacted.	runtime_verified

Severity	Model	Category	Evidence	Verification
CRITICAL	GPT-5.3-Codex / GPT-5.3-Code	unauthenticated_endpoint	The report says <code>`/v1/chat/completions`</code> is mounted publicly, is not gated by <code>`verifyVapiRequest`</code> , and returned HTTP 200 without authentication.	runtime_verified
CRITICAL	GPT-5.3-Codex / GPT-5.3-Code	identity_spoofing	The report says caller number is trusted from the request body and spoofing owner identity reaches owner-aware behavior.	runtime_verified
HIGH	Gemini 3.1 Pro (Preview)	route_registration	The report says mounting <code>`/v1`</code> globally in <code>`index.ts`</code> bypasses standard Express route protection.	source_code_verified
HIGH	DeepSeek V4 Pro (thinking)	identity_spoofing	The report says caller identity and owner context are spoofable via request body phone metadata. Personal identifier redacted.	runtime_verified
HIGH	DeepSeek V4 Pro (thinking)	secret_management	The report says secrets are exposed as plaintext Cloud Run environment variables. Secret values redacted.	runtime_verified
HIGH	Claude Opus 4.8 (fast mode) (Preview)	cost_abuse	The report identifies bill-shock and abuse risk from open paid inference endpoints.	runtime_verified
HIGH	Claude Opus 4.8 (fast mode) (Preview)	route_auth_coverage	The report says <code>`verifyVapiRequest`</code> protects only <code>`/server-events`</code> , not the LLM path.	source_code_verified
HIGH	GPT-5.5	route_auth_coverage	The report says <code>`verifyVapiRequest`</code> does not execute for LLM routes and <code>`/v1`</code> bypasses blanket API auth.	source_code_verified
HIGH	GPT-5.5	owner_context_mapping	The report says the claimed owner mapping fix is not live and deployed source still maps owner to user ID 1.	source_code_verified

Cost And Performance

The benchmark reports did not consistently include token counts, detailed runtime durations, exact retries, approximate dollar cost, or failure-rate percentages. Those values are therefore reported as Not Available rather than inferred.

Table 6. Cost and performance fields. Missing values are explicitly marked Not Available.

Model / Validation	Runtime	Tokens	Approx. Cost	Completion Status	Retries	Failure Rate
GPT-5.5	Not Available	Not Available	Not Available	Completed	Not Available	Not Available
GPT-5 Codex	Not Available	Not Available	Not Available	Completed	Not Available	Not Available
Claude Opus 4.8 Fast	Not Available	Not Available	Not Available	Completed	Not Available	Not Available
DeepSeek V4 Pro Thinking	Not Available	Not Available	Not Available	Completed	Not Available	Not Available
Gemini 3.1 Pro	Not Available	Not Available	Not Available	Completed	Not Available	Not Available
Claude Fable 5	Not Available	Not Available	Not Available	Incomplete	Repeated recovery attempts reported	Not Available
Wingman	3:00+ call segment noted	Not Available	Not Available	Completed validation	Not Available	Qualitative hard fail

Model Chapters

GPT-5.5

Status: extracted sanitized. Verdict: STILL BLOCKED.

[Runtime Verified] runtime: The report says live Cloud Run revision `anara-master-00090-8z2` is latest ready and serving 100% traffic.

[Runtime Verified] security: The report says `POST /v1/chat/completions` returned HTTP 200 from production without `x-vapi-secret`.

[Source-Code Verified] security: The report says `/api/vapi/llm` shares the same handler and lacks `verifyVapiRequest`.

[Runtime Verified] security: The report says spoofing the owner phone number reached owner-aware behavior. Personal identifier redacted.

[Source-Code Verified] security: The report says `v1` mounts `vapiRouter` outside API auth middleware and `verifyVapiRequest` runs for `/server-events`, not LLM routes.

[Source-Code Verified] deployment: The report says 44 dirty files remain and important Vapi fixes are not committed or deployed.

[Source-Code Verified] deployment: The report says deployed HEAD still hardcodes `ownerUserId = isOwner ? 1`, while newer owner-mapping work remains in the working tree.

GPT-5.3-Codex / GPT-5.3-Code

Status: extracted sanitized. Verdict: STILL BLOCKED.

[Runtime Verified] provider_runtime: The report says OpenAI returned `invalid\_api\_key`, while Hugging Face and GitHub Models are live and working.

[Runtime Verified] observability: The report says `/api/vapi/health` returned provider flags and endpoint mapping.

[Runtime Verified] runtime: The report says `anara-master-00090-8z2` is latest ready and serving 100%.

[Runtime Verified] security: The report says unauthenticated POST to `/v1/chat/completions` returned HTTP 200 with a completion.

[Runtime Verified] security: The report says caller number is trusted from the request body and a spoofed owner-number request returned an owner-aware response.

[Source-Code Verified] deployment: The report says ``git status`` and ``git diff`` showed many uncommitted server changes while the service is pinned to a revision/image digest.

[Runtime Verified] reliability: The report says active-revision logs showed no recent ``provider_chain_exhausted``, but only Hugging Face and GitHub are effectively available so exhaustion remains possible if both fail.

Claude Opus 4.8 (fast mode) (Preview)

Status: extracted sanitized. Verdict: STILL BLOCKED.

[Runtime Verified] runtime: The report audited Cloud Run revision ``anara-master-00090-8z2``, serving 100% traffic in project ``anara-voice``, region ``us-central1``.

[Runtime Verified] security: The report says probing ``/v1/chat/completions`` with no secret returned HTTP 200 and a valid AI answer.

[Runtime Verified] security: The report says a spoofed owner phone number caused HTTP 200 and owner-path handling. Personal identifier redacted.

[Source-Code Verified] security: The report says ``/v1`` is mounted outside ``api``, so blanket API auth middleware does not protect it, and ``/api/vapi/llm`` is in ``PUBLIC_API_ROUTES``.

[Source-Code Verified] deployment: The report says ``HEAD=efa83da`` has 44 modified or untracked files and that owner mapping, end-call header, and dedup fixes are uncommitted.

[Source-Code Verified] deployment: The report says deployed ``vapi.ts`` hardcodes ``ownerUserId = isOwner ? 1`` and ignores ``VAPI_OWNER_USER_ID=2``.

[Source-Code Verified] provider_runtime: The report identifies provider chain order as Hugging Face, GitHub, Azure, OpenAI, Gemini.

DeepSeek V4 Pro (thinking)

Status: extracted sanitized. Verdict: STILL BLOCKED.

[Runtime Verified] security: The report says a direct probe to ``/v1/chat/completions`` returned HTTP 200 without any secret header.

[Source-Code Verified] security: The report says ``/api/vapi/llm`` has the same vulnerability and is also in ``PUBLIC_API_ROUTES``.

[Runtime Verified] security: The report says caller identity and owner context are spoofable through ``call.customer.number`` in the request body. Personal identifier redacted.

[Source-Code Verified] deployment: The report says 44 uncommitted files include 676 lines of server changes not deployed to production.

[Source-Code Verified] deployment: The report says deployed code hardcodes ``ownerUserId=1`` while ``VAPI_OWNER_USER_ID=2`` is configured.

[Runtime Verified] provider_runtime: The report says a direct OpenAI ``/v1/models`` probe returned an incorrect API key error.

[Runtime Verified] provider_runtime: The report says Hugging Face and GitHub Models are live and valid, while OpenAI, Azure, and Gemini are unavailable or unconfigured.

Gemini 3.1 Pro (Preview)

Status: extracted sanitized. Verdict: STILL BLOCKED.

[Runtime Verified] runtime: The audit reports that `gcloud run services describe anara-master` verified revision `anara-master-00090-8z2` serving 100% of traffic.

[Runtime Verified] configuration: The audit reports OpenAI, GitHub, and HuggingFace keys present in Cloud Run environment variables, with Azure and Gemini keys missing.

[Runtime Verified] provider_runtime: The audit reports an OpenAI `/v1/models` probe using the production key reference returned HTTP 401 incorrect API key.

[Runtime Verified] provider_runtime: The audit reports HuggingFace `whoami-v2` and GitHub Models `chat/completions` probes returned HTTP 200.

[Source-Code Verified] security: The audit reports `index.ts` and `vapi.ts` show `/v1` mounted outside global API auth middleware and `/api/vapi/llm` whitelisted in `PUBLIC\_API\_ROUTES`.

[Runtime Verified] security: The audit reports a request to `/v1/chat/completions` with no secret header returned HTTP 200 and a completion.

[Runtime Verified] security: The audit reports an owner phone number injected through `call.customer.number` caused HTTP 200 and owner mapping confirmation. Personal identifier redacted.

Claude Fable 5

Status: incomplete session instability. Verdict: INCOMPLETE.

Ranking treatment: exclude from final comparative rankings.

[Manual Verification] benchmark_completion: The report says the forensic investigation could not be completed because the model failed before producing a complete report.

[Manual Verification] session_stability: The report documents repeated `Sorry, no response was returned` failures and a later `400 invalid\_request\_error` involving thinking or redacted_thinking blocks.

[Manual Verification] model_behavior: The report says strengths observed before failure included strong reasoning, good architectural direction, and a promising investigative approach.

[Manual Verification] methodology: The report recommends excluding Fable 5 from final comparative rankings because it never produced a complete benchmark report.

Wingman Validation

Status: extracted sanitized. Verdict: NOT READY FOR PRODUCTION.

[Manual Verification] runtime_stability: The report says task-oriented questions repeatedly triggered the message: `Let us not lose the thread. Something stalled while I was checking that response path.`
Timestamps: 00:32, 00:47, 01:06, 01:13, 01:21, 01:31, 01:46.

[Manual Verification] tool_usage: At 00:59, a live weather request for Albany, NY failed: Anara said she could not pull live data and returned to the stall loop.

[Manual Verification] tool_usage: The category table says scheduling tools triggered crashes during the natural conversation audit.

[Manual Verification] call_termination: At 03:00+, the call entered a two-minute `Take care / Talk soon` repetition loop and could not hang up.

[Manual Verification] memory: The category table says memory failed with zero recall of conversation history due to stalls.

[Manual Verification] voice_experience: The report marks personality, voice quality, turn-taking, latency, emotional consistency, and identity consistency as passing despite backend failures.

Limitations

- Single production system: results reflect ANARA, not every production AI architecture.
- Temporal scope: findings are tied to the July 2, 2026 benchmark environment and reported active revision.
- Provider availability may change after the benchmark window.
- The benchmark evaluates production reasoning, deployment interpretation, and evidence handling rather than raw coding ability alone.
- Reports were supplied as benchmark artifacts; this publication build did not independently re-run every production probe.
- DOCX visual rendering could not be performed in this environment because LibreOffice/soffice was unavailable; PDF rendering was visually sampled.

Recommended Engineering Workflow

1. Use GPT-5.5 for broad initial issue discovery and cross-system framing.
1. Use GPT-5 Codex for repository-aware implementation planning and verification.
1. Use Claude Opus 4.8 Fast for independent forensic challenge.
1. Use DeepSeek V4 Pro Thinking for adversarial risk expansion.
1. Use Gemini 3.1 Pro for independent cross-checking.
1. Use Wingman as live acceptance testing, not as a peer model score.
1. Promote to production only after security, deployment, runtime, and live-call gates pass.

Appendices

Complete Extracted Findings

Table 7. Complete extracted findings.

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	runtime	The audit reports that <code>`gcloud run services describe anara-master`</code> verified revision <code>`anara-master-00090-8z2`</code> serving 100% of traffic.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 1
Gemini 3.1 Pro (Preview)	configuration	The audit reports OpenAI, GitHub, and HuggingFace keys present in Cloud Run environment variables, with Azure and Gemini keys missing.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 2
Gemini 3.1 Pro (Preview)	provider_runtime	The audit reports an OpenAI <code>`/v1/models`</code> probe using the production key reference returned HTTP 401 incorrect API key.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 3

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	provider_runtime	The audit reports HuggingFace `whoami-v2` and GitHub Models `chat/completions` probes returned HTTP 200.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 4
Gemini 3.1 Pro (Preview)	security	The audit reports `index.ts` and `vapi.ts` show `/v1` mounted outside global API auth middleware and `/api/vapi/llm` whitelisted in `PUBLIC_API_ROUTES`.	[Source-Code Verified]	benchmark-data/gemini.md: Timeline of Findings, item 5
Gemini 3.1 Pro (Preview)	security	The audit reports a request to `/v1/chat/completions` with no secret header returned HTTP 200 and a completion.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 6
Gemini 3.1 Pro (Preview)	security	The audit reports an owner phone number injected through `call.customer.number` caused HTTP 200 and owner mapping confirmation. Personal identifier redacted.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 7
Gemini 3.1 Pro (Preview)	deployment	The audit reports `git diff HEAD server/routes/vapi.ts` showed `resolveOwnerInternalUserId()` and `x-anara-end-call` changes only in uncommitted files.	[Source-Code Verified]	benchmark-data/gemini.md: Timeline of Findings, item 8
Gemini 3.1 Pro (Preview)	production_correctness	The audit reports `resolveNumericUserIdForLookups` prevents phone-number strings or prefixes from being parsed into Postgres integer IDs.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right
Gemini 3.1 Pro (Preview)	production_runtime	The audit reports weather succeeds and `OPENWEATHER_API_KEY` is configured correctly.	[Runtime Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right
Gemini 3.1 Pro (Preview)	production_resilience	The audit reports Calendar and Gmail failures are isolated by try/catch buffers and produce fallback strings instead of crashing the chat path.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right
Gemini 3.1 Pro (Preview)	observability	The audit reports `getLlmProviderDiagnostics()` startup logging is present.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	deployment	The audit reports 44 dirty files in the working directory and says testing is detached from deployed production behavior.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Missed; Hidden Risks
Gemini 3.1 Pro (Preview)	deployment	The audit reports `VAPI_OWNER_USER_ID=2` in environment, while deployed code hardcodes owner user ID `1`.	[Source-Code Verified]	benchmark-data/gemini.md: Hidden Risks
Gemini 3.1 Pro (Preview)	reliability	The audit says Azure and Gemini are disconnected, OpenAI has a dead key, and the system relies entirely on HuggingFace and GitHub Models.	[AI Interpretation]	benchmark-data/gemini.md: What Previous Engineering Missed
Gemini 3.1 Pro (Preview)	runtime_stability	The audit says goodbye regex phrases such as `i'm good` and `take care` risk accidental call termination.	[AI Interpretation]	benchmark-data/gemini.md: Hidden Risks
DeepSeek V4 Pro (thinking)	security	The report says a direct probe to `/v1/chat/completions` returned HTTP 200 without any secret header.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	security	The report says `/api/vapi/llm` has the same vulnerability and is also in `PUBLIC_API_ROUTES`.	[Source-Code Verified]	benchmark-data/deepseek.md: Security Findings
DeepSeek V4 Pro (thinking)	security	The report says caller identity and owner context are spoofable through `call.customer.number` in the request body. Personal identifier redacted.	[Runtime Verified]	benchmark-data/deepseek.md: Security Findings
DeepSeek V4 Pro (thinking)	deployment	The report says 44 uncommitted files include 676 lines of server changes not deployed to production.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	deployment	The report says deployed code hardcodes `ownerUserId=1` while `VAPI_OWNER_USER_ID=2` is configured.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	provider_runtime	The report says a direct OpenAI `/v1/models` probe returned an incorrect API key error.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings; Evidence Appendix
DeepSeek V4 Pro (thinking)	provider_runtime	The report says Hugging Face and GitHub Models are live and valid, while OpenAI, Azure, and Gemini are unavailable or unconfigured.	[Runtime Verified]	benchmark-data/deepseek.md: Evidence Appendix

Source	Category	Evidence	Verification	Location
DeepSeek V4 Pro (thinking)	security	The report says Cloud Run revision metadata exposes keys and secrets in plaintext environment variables. Secret values redacted.	[Runtime Verified]	benchmark-data/deepseek.md: Security Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	runtime_stability	The report says the deployed max-turn guard lacks the `x-anara-end-call` header and that the `maxTurnClosedCalls` map only exists in uncommitted code.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings; Remaining Production Blockers
DeepSeek V4 Pro (thinking)	production_resilience	The report confirms weather works and calendar/Gmail degrade gracefully through independent try/catch patterns.	[Source-Code Verified]	benchmark-data/deepseek.md: What Previous Engineering Got Right
DeepSeek V4 Pro (thinking)	runtime_stability	The report notes `maxTurnClosedCalls` and circuit breaker state are in-memory, so behavior does not persist across Cloud Run instances or restarts.	[AI Interpretation]	benchmark-data/deepseek.md: Hidden Risks
Claude Opus 4.8 (fast mode) (Preview)	runtime	The report audited Cloud Run revision `anara-master-00090-8z2`, serving 100% traffic in project `anara-voice`, region `us-central1`.	[Runtime Verified]	benchmark-data/opus48.md: Audit Metadata; Timeline of Findings
Claude Opus 4.8 (fast mode) (Preview)	security	The report says probing `/v1/chat/completions` with no secret returned HTTP 200 and a valid AI answer.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T8; Security Findings
Claude Opus 4.8 (fast mode) (Preview)	security	The report says a spoofed owner phone number caused HTTP 200 and owner-path handling. Personal identifier redacted.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T9; Security Findings
Claude Opus 4.8 (fast mode) (Preview)	security	The report says `/v1` is mounted outside `api`, so blanket API auth middleware does not protect it, and `/api/vapi/llm` is in `PUBLIC_API_ROUTES`.	[Source-Code Verified]	benchmark-data/opus48.md: What Previous Engineering Missed
Claude Opus 4.8 (fast mode) (Preview)	deployment	The report says `HEAD=efa83da` has 44 modified or untracked files and that owner mapping, end-call header, and dedup fixes are uncommitted.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T1 and T11
Claude Opus 4.8 (fast mode) (Preview)	deployment	The report says deployed `vapi.ts` hardcodes `ownerUserId = isOwner ? 1` and ignores `VAPI_OWNER_USER_ID=2`.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T10

Source	Category	Evidence	Verification	Location
Claude Opus 4.8 (fast mode) (Preview)	provider_runtime	The report identifies provider chain order as Hugging Face, GitHub, Azure, OpenAI, Gemini.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T6
Claude Opus 4.8 (fast mode) (Preview)	provider_runtime	The report says OpenAI returned HTTP 401 while Hugging Face and GitHub Models both returned HTTP 200.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T4-T5
Claude Opus 4.8 (fast mode) (Preview)	security	The report says Cloud Run revision metadata exposes credentials as plaintext environment variables. Secret values redacted.	[Runtime Verified]	benchmark-data/opus48.md: What Previous Engineering Missed; Security Findings
Claude Opus 4.8 (fast mode) (Preview)	production_resilience	The report verifies weather, Calendar, Gmail, provider diagnostics, and phone user-id safety as working or gracefully degrading.	[Source-Code Verified]	benchmark-data/opus48.md: What Previous Engineering Got Right
Claude Opus 4.8 (fast mode) (Preview)	runtime_stability	The report says `x-anara-end-call` and repeat deduplication are only in uncommitted code, leaving deployed call termination weaker than claimed.	[Source-Code Verified]	benchmark-data/opus48.md: What Previous Engineering Missed; Remaining Production Blockers
GPT-5.5	runtime	The report says live Cloud Run revision `anara-master-00090-8z2` is latest ready and serving 100% traffic.	[Runtime Verified]	benchmark-data/gpt55.md: Executive Summary; Timeline of Findings
GPT-5.5	security	The report says `POST /v1/chat/completions` returned HTTP 200 from production without `x-vapi-secret`.	[Runtime Verified]	benchmark-data/gpt55.md: Executive Summary; Timeline of Findings
GPT-5.5	security	The report says `/api/vapi/llm` shares the same handler and lacks `verifyVapiRequest`.	[Source-Code Verified]	benchmark-data/gpt55.md: What Previous Engineering Missed
GPT-5.5	security	The report says spoofing the owner phone number reached owner-aware behavior. Personal identifier redacted.	[Runtime Verified]	benchmark-data/gpt55.md: Executive Summary; Security Findings
GPT-5.5	security	The report says `/v1` mounts `vapiRouter` outside API auth middleware and `verifyVapiRequest` runs for `/server-events`, not LLM routes.	[Source-Code Verified]	benchmark-data/gpt55.md: Timeline of Findings; Security Findings
GPT-5.5	deployment	The report says 44 dirty files remain and important Vapi fixes are not committed or deployed.	[Source-Code Verified]	benchmark-data/gpt55.md: Timeline of Findings; What Previous Engineering Missed

Source	Category	Evidence	Verification	Location
GPT-5.5	deployment	The report says deployed HEAD still hardcodes `ownerUserId = isOwner ? 1`, while newer owner-mapping work remains in the working tree.	[Source-Code Verified]	benchmark-data/gpt55.m d: What Previous Engineering Missed
GPT-5.5	runtime_stability	The report says live code lacks the uncommitted `x-anara-end-call` header and repeat-close cache.	[Source-Code Verified]	benchmark-data/gpt55.m d: What Previous Engineering Missed
GPT-5.5	provider_runtime	The report says OpenAI returned HTTP 401, while GitHub Models and Hugging Face probes returned HTTP 200.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	production_resilience	The report says weather succeeds, Gmail and Calendar degrade gracefully, and provider diagnostics are present.	[Source-Code Verified]	benchmark-data/gpt55.m d: What Previous Engineering Got Right
GPT-5.5	security	The report says Cloud Run revision exposes secret values as direct environment variables rather than Secret Manager bindings. Values redacted.	[Runtime Verified]	benchmark-data/gpt55.m d: What Previous Engineering Missed; Security Findings
GPT-5.3-Codex / GPT-5.3-Code	provider_runtime	The report says OpenAI returned `invalid_api_key`, while Hugging Face and GitHub Models are live and working.	[Runtime Verified]	benchmark-data/codex.m d: Executive Summary; What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	observability	The report says `/api/vapi/health` returned provider flags and endpoint mapping.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	runtime	The report says `anara-master-00090-8z2` is latest ready and serving 100%.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	security	The report says unauthenticated POST to `/v1/chat/completions` returned HTTP 200 with a completion.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	security	The report says caller number is trusted from the request body and a spoofed owner-number request returned an owner-aware response.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	deployment	The report says `git status` and `git diff` showed many uncommitted server changes while the service is pinned to a revision/image digest.	[Source-Code Verified]	benchmark-data/codex.m d: What Codex Missed

Source	Category	Evidence	Verification	Location
GPT-5.3-Codex / GPT-5.3-Code	reliability	The report says active-revision logs showed no recent `provider_chain_exhausted`, but only Hugging Face and GitHub are effectively available so exhaustion remains possible if both fail.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	production_resilience	The report says phone-number-to-integer-ID protection is largely correct and Gmail/Calendar degradation is isolated with try/catch and fallback notes.	[Source-Code Verified]	benchmark-data/codex.m d: What Codex Got Right
Wingman Validation	runtime_stability	The report says task-oriented questions repeatedly triggered the message: `Let us not lose the thread. Something stalled while I was checking that response path.` Timestamps: 00:32, 00:47, 01:06, 01:13, 01:21, 01:31, 01:46.	[Manual Verification]	benchmark-data/wingma n.md: Production-Blocking Issues
Wingman Validation	tool_usage	At 00:59, a live weather request for Albany, NY failed: Anara said she could not pull live data and returned to the stall loop.	[Manual Verification]	benchmark-data/wingma n.md: Production-Blocking Issues
Wingman Validation	tool_usage	The category table says scheduling tools triggered crashes during the natural conversation audit.	[Manual Verification]	benchmark-data/wingma n.md: Category Results
Wingman Validation	call_termination	At 03:00+, the call entered a two-minute `Take care / Talk soon` repetition loop and could not hang up.	[Manual Verification]	benchmark-data/wingma n.md: Production-Blocking Issues
Wingman Validation	memory	The category table says memory failed with zero recall of conversation history due to stalls.	[Manual Verification]	benchmark-data/wingma n.md: Category Results
Wingman Validation	voice_experience	The report marks personality, voice quality, turn-taking, latency, emotional consistency, and identity consistency as passing despite backend failures.	[Manual Verification]	benchmark-data/wingma n.md: Category Results
Claude Fable 5	benchmark_completion	The report says the forensic investigation could not be completed because the model failed before producing a complete report.	[Manual Verification]	benchmark-data/fable.md : Investigation Status; Benchmark Result

Source	Category	Evidence	Verification	Location
Claude Fable 5	session_stability	The report documents repeated `Sorry, no response was returned` failures and a later `400 invalid_request_error` involving thinking or redacted_thinking blocks.	[Manual Verification]	benchmark-data/fable.md : Observed Errors
Claude Fable 5	model_behavior	The report says strengths observed before failure included strong reasoning, good architectural direction, and a promising investigative approach.	[Manual Verification]	benchmark-data/fable.md : Stability Assessment
Claude Fable 5	methodology	The report recommends excluding Fable 5 from final comparative rankings because it never produced a complete benchmark report.	[Manual Verification]	benchmark-data/fable.md : Methodology Treatment

Runtime Evidence

Table 8. Runtime observations.

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	cloud_run_runtime	Cloud Run revision `anar a-master-00090-8z2` reportedly serves 100% of traffic.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 1
Gemini 3.1 Pro (Preview)	endpoint_runtime	Unauthenticated runtime probe to `/v1/chat/completions` reportedly returned HTTP 200 and a completion.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 6
Gemini 3.1 Pro (Preview)	endpoint_runtime	Owner spoofing runtime probe reportedly returned HTTP 200 and owner mapping confirmation. Personal identifier redacted.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 7
DeepSeek V4 Pro (thinking)	endpoint_runtime	The report says a direct runtime probe to `/v1/chat/completions` returned HTTP 200 without a secret header.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings; Evidence Appendix
DeepSeek V4 Pro (thinking)	provider_runtime	The report says Hugging Face and GitHub Models are working providers, while OpenAI returns 401 and Azure/Gemini are not configured.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings; Evidence Appendix
DeepSeek V4 Pro (thinking)	phone_pipeline	The report says the phone pipeline is functional and the system can take Wingman calls.	[Runtime Verified]	benchmark-data/deepseek.md: Executive Summary
Claude Opus 4.8 (fast mode) (Preview)	cloud_run_runtime	The report audited live revision `anara-master-00090-8z2`, serving 100% traffic.	[Runtime Verified]	benchmark-data/opus48.md: Audit Metadata

Source	Category	Evidence	Verification	Location
Claude Opus 4.8 (fast mode) (Preview)	endpoint_runtime	The report says <code>`/v1/chat/completions`</code> returned HTTP 200 and a valid AI answer with no secret header.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T8
Claude Opus 4.8 (fast mode) (Preview)	endpoint_runtime	The report says spoofed owner metadata returned HTTP 200 and owner-path handling. Personal identifier redacted.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T9
Claude Opus 4.8 (fast mode) (Preview)	phone_pipeline	The report says the phone pipeline functions and two providers, Hugging Face plus GitHub Models, are live.	[Runtime Verified]	benchmark-data/opus48.md: Executive Summary
GPT-5.5	cloud_run_runtime	The report says live Cloud Run revision <code>`anar a-master-00090-8z2`</code> is serving 100% traffic.	[Runtime Verified]	benchmark-data/gpt55.md: Executive Summary; Timeline of Findings
GPT-5.5	endpoint_runtime	The report says <code>`/v1/chat/completions`</code> returned HTTP 200 from production without <code>`x-vapi-secret`</code> .	[Runtime Verified]	benchmark-data/gpt55.md: Timeline of Findings
GPT-5.5	endpoint_runtime	The report says spoofed owner metadata returned HTTP 200 and reached owner-number-aware behavior. Personal identifier redacted.	[Runtime Verified]	benchmark-data/gpt55.md: Timeline of Findings
GPT-5.5	provider_runtime	The report says Hugging Face and GitHub probes returned HTTP 200, while OpenAI returned HTTP 401.	[Runtime Verified]	benchmark-data/gpt55.md: Timeline of Findings
GPT-5.3-Codex / GPT-5.3-Code	cloud_run_runtime	The report says <code>`anara-master-00090-8z2`</code> is latest ready and serving 100%.	[Runtime Verified]	benchmark-data/codex.md: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	health_endpoint	The report says <code>`/api/vapi/health`</code> returned provider flags and endpoint mapping.	[Runtime Verified]	benchmark-data/codex.md: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	endpoint_runtime	The report says unauthenticated POST to <code>`/v1/chat/completions`</code> returned HTTP 200 with a completion.	[Runtime Verified]	benchmark-data/codex.md: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	endpoint_runtime	The report says a spoofed owner-number request returned an owner-aware response.	[Runtime Verified]	benchmark-data/codex.md: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	provider_runtime	The report says active-revision log search showed no recent <code>`provider_chain_exhausted`</code> events.	[Runtime Verified]	benchmark-data/codex.md: What Codex Missed

Source	Category	Evidence	Verification	Location
Wingman Validation	conversation_runtime	Natural conversation audit observed repeated response-path stall messages during task-oriented requests.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	tool_runtime	Weather and scheduling tool attempts triggered crashes or stall-loop behavior.	[Manual Verification]	benchmark-data/wingman.md: Category Results; Production-Blocking Issues
Wingman Validation	call_termination	The call ended in a two-minute polite repetition loop and could not hang up cleanly.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	latency	Average latency was reported as 1.5 seconds and acceptable for voice, but responses were broken.	[Manual Verification]	benchmark-data/wingman.md: Category Results
Claude Fable 5	model_runtime_stability	The benchmark session failed before the model produced a complete forensic report.	[Manual Verification]	benchmark-data/fable.md : Investigation Status; Observed Errors
Claude Fable 5	model_runtime_stability	Subsequent attempts returned client-side errors, including no-response failures and a 400 invalid_request_error related to thinking/redacted_thinking blocks.	[Manual Verification]	benchmark-data/fable.md : Observed Errors

Production Findings

Table 9. Production findings.

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	production_readiness	Final verdict is STILL BLOCKED; production readiness score is 4/10.	[AI Interpretation]	benchmark-data/gemini.md: Production Readiness Score; Final Verdict
Gemini 3.1 Pro (Preview)	subsystem_health	Weather succeeds with `OPENWEATHER_API_KEY` configured.	[Runtime Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right
Gemini 3.1 Pro (Preview)	subsystem_degradation	Calendar and Gmail degrade gracefully through isolated error handling and fallback strings.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Got Right
Gemini 3.1 Pro (Preview)	provider_resilience	Provider chain remains fragile because Azure and Gemini are disconnected, OpenAI is invalid, and runtime depends on HuggingFace plus GitHub Models.	[AI Interpretation]	benchmark-data/gemini.md: What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	production_readiness	Final verdict is STILL BLOCKED; production readiness score is 5/10.	[AI Interpretation]	benchmark-data/deepseek.md: Production Readiness Score; Final Verdict

Source	Category	Evidence	Verification	Location
DeepSeek V4 Pro (thinking)	production_readiness	The report scores LLM availability 7/10 and phone pipeline 7/10, noting two working providers and functional calls but owner mapping issues.	[AI Interpretation]	benchmark-data/deepseek.md: Production Readiness Score
DeepSeek V4 Pro (thinking)	security_readiness	The report scores security 2/10 due to critical auth bypass and plaintext secrets.	[AI Interpretation]	benchmark-data/deepseek.md: Production Readiness Score
DeepSeek V4 Pro (thinking)	subsystem_degradation	The report confirms graceful degradation for weather, calendar, and Gmail paths.	[Source-Code Verified]	benchmark-data/deepseek.md: What Previous Engineering Got Right
Claude Opus 4.8 (fast mode) (Preview)	production_readiness	Final verdict is STILL BLOCKED; production readiness score is 5/10.	[AI Interpretation]	benchmark-data/opus48.md: Production Readiness Score; Final Verdict
Claude Opus 4.8 (fast mode) (Preview)	production_readiness	The report scores LLM availability 7/10, phone pipeline correctness 6/10, security 2/10, deployment integrity 4/10, graceful degradation 8/10, and observability 7/10.	[AI Interpretation]	benchmark-data/opus48.md: Production Readiness Score
Claude Opus 4.8 (fast mode) (Preview)	subsystem_degradation	The report says weather, Gmail/Calendar/Digital Twin degradation, phone user-id safety, and provider diagnostics are verified working.	[Source-Code Verified]	benchmark-data/opus48.md: Final Verdict; What Previous Engineering Got Right
Claude Opus 4.8 (fast mode) (Preview)	production_prioritization	The report says OpenAI rotation, Secret Manager migration, and adding a third provider are high-priority but not blockers for the next call.	[AI Interpretation]	benchmark-data/opus48.md: Final Verdict
GPT-5.5	production_readiness	Final verdict is STILL BLOCKED; production readiness score is 5/10.	[AI Interpretation]	benchmark-data/gpt55.md: Production Readiness Score; Final Verdict
GPT-5.5	production_readiness	The report says the app can answer and core degradation paths are improved, but unauthenticated LLM endpoint and deployment drift prevent a production call.	[AI Interpretation]	benchmark-data/gpt55.md: Production Readiness Score
GPT-5.5	subsystem_degradation	The report verifies weather, Gmail, Calendar, and provider diagnostics as working or gracefully degrading.	[Source-Code Verified]	benchmark-data/gpt55.md: What Previous Engineering Got Right
GPT-5.5	production_prioritization	The report says OpenAI 401 is not a Wingman blocker by itself because Hugging Face and GitHub currently answer successfully.	[AI Interpretation]	benchmark-data/gpt55.md: Remaining Production Blockers

Source	Category	Evidence	Verification	Location
GPT-5.3-Codex / GPT-5.3-Code	production_readiness	Final verdict is STILL BLOCKED; production readiness score is 4/10.	[AI Interpretation]	benchmark-data/codex.m d: Production Readiness Score; Final Verdict
GPT-5.3-Codex / GPT-5.3-Code	production_readiness	The report says production is not ready for another Wingman call because of unauthenticated endpoint access, spoofable caller identity, and deployment drift.	[AI Interpretation]	benchmark-data/codex.m d: Executive Summary; Remaining Production Blockers
GPT-5.3-Codex / GPT-5.3-Code	subsystem_degradation	The report says phone-number ID protection is largely correct and Gmail/Calendar degradation uses isolated try/catch fallback notes.	[Source-Code Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	production_prioritization	The report says OpenAI 401 alone is not the blocker because Hugging Face and GitHub are currently working.	[AI Interpretation]	benchmark-data/codex.m d: Remaining Production Blockers
Wingman Validation	production_readiness	Final verdict is NOT READY FOR PRODUCTION; overall score is 2/10.	[Manual Verification]	benchmark-data/wingma n.md: Executive Summary; Final Verdict
Wingman Validation	production_blocker	The report says Anara was non-functional for back-office tasks because task-oriented questions triggered response-path stalls.	[Manual Verification]	benchmark-data/wingma n.md: Executive Summary; Production-Blocking Issues
Wingman Validation	voice_experience	Voice quality, initial personality, turn-taking, latency, emotional consistency, and identity consistency passed.	[Manual Verification]	benchmark-data/wingma n.md: Category Results
Claude Fable 5	benchmark_completion	No ANARA production readiness findings could be independently verified because the investigation terminated before reaching a conclusion.	[Manual Verification]	benchmark-data/fable.md : Benchmark Result
Claude Fable 5	production_readiness	Final verdict is INCOMPLETE; no production readiness score was assigned.	[Manual Verification]	benchmark-data/fable.md : Final Verdict

Deployment Findings

Table 10. Deployment findings.

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	cloud_run_revision	The live Cloud Run revision is reported as `a nara-master-00090-8z2`, serving 100% of traffic.	[Runtime Verified]	benchmark-data/gemini. md: Timeline of Findings, item 1

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	deployment_drift	The report says 44 files are dirty and critical VAPI fixes are uncommitted.	[Source-Code Verified]	benchmark-data/gemini.md: What Previous Engineering Missed; Hidden Risks
Gemini 3.1 Pro (Preview)	configuration_drift	`VAPI_OWNER_USER_ID=2` is reportedly configured, but deployed code hardcodes owner user ID `1`.	[Source-Code Verified]	benchmark-data/gemini.md: Hidden Risks
Gemini 3.1 Pro (Preview)	release_blocker	The report says owner ID resolution and `x-anara-end-call` state mapping must be tested, committed, and deployed.	[AI Interpretation]	benchmark-data/gemini.md: Remaining Production Blockers
DeepSeek V4 Pro (thinking)	deployment_drift	The report says 44 uncommitted files and 676 added server lines are not deployed.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	configuration_drift	The report says deployed code hardcodes `ownerUserId=1` while `VAPI_OWNER_USER_ID=2` is configured.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings; What Previous Engineering Missed
DeepSeek V4 Pro (thinking)	release_blocker	The report says max-turn `x-anara-end-call` and owner ID resolution must be committed and deployed before the next Wingman call.	[AI Interpretation]	benchmark-data/deepseek.md: Remaining Production Blockers; Final Verdict
Claude Opus 4.8 (fast mode) (Preview)	cloud_run_revision	The report says live revision `anara-master-00090-8z2` serves 100% traffic in project `anara-voice`, region `us-central1`.	[Runtime Verified]	benchmark-data/opus48.md: Audit Metadata
Claude Opus 4.8 (fast mode) (Preview)	deployment_drift	The report says `HEAD=efa83da` has 44 modified or untracked files and the claimed owner-mapping fix is not deployed.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T1 and T11
Claude Opus 4.8 (fast mode) (Preview)	configuration_drift	The report says deployed code hardcodes owner user ID `1` while `VAPI_OWNER_USER_ID=2` is configured.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T3 and T10
Claude Opus 4.8 (fast mode) (Preview)	release_blocker	The report says call-termination hardening, including `x-anara-end-call` and repeat deduplication, is uncommitted and not deployed.	[Source-Code Verified]	benchmark-data/opus48.md: Remaining Production Blockers
GPT-5.5	cloud_run_revision	The report says live revision `anara-master-00090-8z2` is latest ready and serving 100% traffic.	[Runtime Verified]	benchmark-data/gpt55.md: Timeline of Findings
GPT-5.5	deployment_drift	The report says 44 dirty files include uncommitted Vapi fixes that are not deployed.	[Source-Code Verified]	benchmark-data/gpt55.md: Timeline of Findings; What Previous Engineering Missed

Source	Category	Evidence	Verification	Location
GPT-5.5	configuration_drift	The report says Cloud Run has `VAPI_OWNER_USER_ID=2`, but deployed code hardcodes owner user ID `1`.	[Source-Code Verified]	benchmark-data/gpt55.md: Hidden Risks
GPT-5.5	release_blocker	The report says owner mapping and max-turn call-termination fixes must be deployed before another Wingman production call.	[AI Interpretation]	benchmark-data/gpt55.md: Remaining Production Blockers
GPT-5.3-Codex / GPT-5.3-Code	cloud_run_revision	The report says `anara-master-00090-8z2` is the current Cloud Run revision serving 100%.	[Runtime Verified]	benchmark-data/codex.md: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	deployment_drift	The report says many uncommitted server changes exist while the service is pinned to a revision/image digest, creating a deployment consistency gap.	[Source-Code Verified]	benchmark-data/codex.md: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	release_blocker	The report says reported VAPI safety fixes are not fully verifiable as deployed behavior.	[AI Interpretation]	benchmark-data/codex.md: Remaining Production Blockers
Wingman Validation	release_blocker	The report says production is not ready because the phone backend response paths fail during task-oriented live conversation.	[Manual Verification]	benchmark-data/wingman.md: Executive Summary; Final Verdict
Wingman Validation	follow_up_required	The report recommends investigating Vapi and LLM logs for the `stalled response path` error.	[Manual Verification]	benchmark-data/wingman.md: Engineering Note

Timeline

Table 11. Full event timeline.

Source	Category	Event	Verification	Location
Gemini 3.1 Pro (Preview)	runtime	Cloud Run audit verified revision `anara-master-00090-8z2` serving all traffic.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 1
Gemini 3.1 Pro (Preview)	configuration	Environment audit found OpenAI, GitHub, and HuggingFace keys present; Azure and Gemini keys missing.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 2
Gemini 3.1 Pro (Preview)	provider_runtime	OpenAI provider probe returned HTTP 401 incorrect API key.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 3
Gemini 3.1 Pro (Preview)	provider_runtime	HuggingFace and GitHub Models provider probes returned HTTP 200.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 4

Source	Category	Event	Verification	Location
Gemini 3.1 Pro (Preview)	security	Code audit found <code>`v1`</code> outside global auth and <code>`api/vapi/llm`</code> public.	[Source-Code Verified]	benchmark-data/gemini.md: Timeline of Findings, item 5
Gemini 3.1 Pro (Preview)	security	Public runtime probe to <code>`v1/chat/completions`</code> without secret header returned HTTP 200.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 6
Gemini 3.1 Pro (Preview)	security	Owner-spoofing probe using injected caller metadata returned HTTP 200 and owner-mapping evidence.	[Runtime Verified]	benchmark-data/gemini.md: Timeline of Findings, item 7
Gemini 3.1 Pro (Preview)	deployment	Deployment drift inspection found owner ID resolution and end-call changes only in uncommitted code.	[Source-Code Verified]	benchmark-data/gemini.md: Timeline of Findings, item 8
DeepSeek V4 Pro (thinking)	security	<code>`v1/chat/completions`</code> direct probe returned HTTP 200 without a secret header.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 1
DeepSeek V4 Pro (thinking)	deployment	44 uncommitted files and 676 server lines were found not deployed to production.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 2
DeepSeek V4 Pro (thinking)	provider_runtime	OpenAI <code>`v1/models`</code> probe returned incorrect API key.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 3
DeepSeek V4 Pro (thinking)	deployment	Deployed code was found hardcoding <code>`ownerUserId=1`</code> while <code>`V_API_OWNER_USER_ID=2`</code> was configured.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 4
DeepSeek V4 Pro (thinking)	configuration	Gemini provider was not configured in Cloud Run environment.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 5
DeepSeek V4 Pro (thinking)	configuration	Azure OpenAI provider was not configured in Cloud Run environment.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 6
DeepSeek V4 Pro (thinking)	runtime_stability	Deployed max-turn guard was found lacking <code>`x-anara-end-call`</code> header.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 7
DeepSeek V4 Pro (thinking)	security	Cloud Run revision metadata reportedly exposed secrets as plaintext environment variables. Values redacted.	[Runtime Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 8
DeepSeek V4 Pro (thinking)	runtime_stability	<code>`maxTurnClosedCalls`</code> map was found only in uncommitted code.	[Source-Code Verified]	benchmark-data/deepseek.md: Timeline of Findings, item 9
Claude Opus 4.8 (fast mode) (Preview)	methodology	Read deployment and debug protocols, establishing deployment flow and drift warning.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T0

Source	Category	Event	Verification	Location
Claude Opus 4.8 (fast mode) (Preview)	deployment	`git log` and `git status` found `HEAD=efa83da` with 44 modified or untracked files.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T1
Claude Opus 4.8 (fast mode) (Preview)	runtime	`gcloud run services describe` found live revision `00090-8z2` serving 100% traffic.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T2
Claude Opus 4.8 (fast mode) (Preview)	configuration	`gcloud run revisions describe` found plaintext environment variables, `V_API_OWNER_USER_ID=2`, and no Gemini/Azure keys.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T3
Claude Opus 4.8 (fast mode) (Preview)	provider_runtime	OpenAI production key probe returned HTTP 401 incorrect API key.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T4
Claude Opus 4.8 (fast mode) (Preview)	provider_runtime	GitHub Models and Hugging Face token probes both returned HTTP 200.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T5
Claude Opus 4.8 (fast mode) (Preview)	provider_runtime	Provider chain in `ai-chat-service.ts` read as Hugging Face, GitHub, Azure, OpenAI, Gemini.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T6
Claude Opus 4.8 (fast mode) (Preview)	observability	Production startup diagnostics showed OpenAI configured, Gemini and Azure not configured.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T7
Claude Opus 4.8 (fast mode) (Preview)	security	Probe to `/v1/chat/completions` with no secret returned HTTP 200 and valid AI answer.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T8
Claude Opus 4.8 (fast mode) (Preview)	security	Probe with spoofed owner number returned HTTP 200 and owner-path handling. Personal identifier redacted.	[Runtime Verified]	benchmark-data/opus48.md: Timeline of Findings, T9
Claude Opus 4.8 (fast mode) (Preview)	deployment	`git show HEAD:server/routes/vapi.ts` found hardcoded owner user ID and ignored `VAPI_OWNER_USER_ID=2`.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T10
Claude Opus 4.8 (fast mode) (Preview)	deployment	Working tree diff found owner mapping, end-call header, and dedup fixes uncommitted and undeployed.	[Source-Code Verified]	benchmark-data/opus48.md: Timeline of Findings, T11
GPT-5.5	runtime	Cloud Run service check found `anara-master-00090-8z2` latest ready and serving 100% traffic.	[Runtime Verified]	benchmark-data/gpt55.md: Timeline of Findings

Source	Category	Event	Verification	Location
GPT-5.5	configuration	Revision environment check found OpenAI, GitHub, and Hugging Face configured; Azure and Gemini not configured.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	provider_runtime	OpenAI `/v1/models` probe returned HTTP 401 incorrect API key.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	provider_runtime	GitHub Models probe returned HTTP 200.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	provider_runtime	Hugging Face `whoami-v2` probe returned HTTP 200.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	security	Auth bypass probe to `/v1/chat/completions` without secret returned HTTP 200 with model response.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	security	Owner spoof probe with spoofed owner phone number returned HTTP 200 and owner-number behavior. Personal identifier redacted.	[Runtime Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	security	Source trace found `/v1` mounts `vapiRouter` outside API auth middleware.	[Source-Code Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	security	Source trace found `verifyVapiRequest` runs for `/server-events`, not LLM routes.	[Source-Code Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.5	deployment	Deployment drift check found 44 dirty files and important Vapi fixes not committed or deployed.	[Source-Code Verified]	benchmark-data/gpt55.m d: Timeline of Findings
GPT-5.3-Codex / GPT-5.3-Code	observability	Live health endpoint confirmed provider flags and endpoint mapping.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	provider_runtime	OpenAI models endpoint probe returned invalid API key.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	runtime	Current Cloud Run revision `anara-master-00090-8z2` confirmed latest ready and serving 100%.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Got Right
GPT-5.3-Codex / GPT-5.3-Code	security	Unauthenticated POST to `/v1/chat/completions` returned HTTP 200 with completion.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	security	Spoofed owner-number request returned owner-aware response.	[Runtime Verified]	benchmark-data/codex.m d: What Codex Missed

Source	Category	Event	Verification	Location
GPT-5.3-Codex / GPT-5.3-Code	reliability	Active-revision log search showed no recent `provider_chain_exhausted`, but architecture still permits exhaustion if Hugging Face and GitHub fail.	[Runtime Verified]	benchmark-data/codex.mmd: What Codex Missed
GPT-5.3-Codex / GPT-5.3-Code	deployment	Deployment drift evidence found many uncommitted server changes while production is pinned to a revision/image digest.	[Source-Code Verified]	benchmark-data/codex.mmd: What Codex Missed
Wingman Validation	conversation_runtime	At 00:32, task-oriented conversation triggered the first reported response-path stall.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 00:47, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	tool_runtime	At 00:59, live weather request for Albany, NY failed and returned to the stall loop.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 01:06, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 01:13, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 01:21, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 01:31, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	conversation_runtime	At 01:46, response-path stall repeated.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Wingman Validation	call_termination	At 03:00+, call entered a two-minute polite hangup loop.	[Manual Verification]	benchmark-data/wingman.md: Production-Blocking Issues
Claude Fable 5	benchmark_execution	Model began an attempted forensic investigation of ANARA production readiness, Cloud Run configuration, request path, deployment configuration, and runtime behavior.	[Manual Verification]	benchmark-data/fable.md: Findings Prior To Failure
Claude Fable 5	benchmark_execution	Investigation failed before completion and no production findings were independently verified.	[Manual Verification]	benchmark-data/fable.md: Investigation Status; Benchmark Result

Source	Category	Event	Verification	Location
Claude Fable 5	model_runtime_stability	Repeated retry attempts returned 'Sorry, no response was returned'.	[Manual Verification]	benchmark-data/fable.md : Observed Errors
Claude Fable 5	model_runtime_stability	Later attempts produced a 400 invalid_request_error related to thinking or redacted_thinking blocks.	[Manual Verification]	benchmark-data/fable.md : Observed Errors
Claude Fable 5	methodology	Benchmark treatment set to INCOMPLETE and excluded from final comparative rankings.	[Manual Verification]	benchmark-data/fable.md : Methodology Treatment; Final Verdict

Cost Observations

Table 12. Cost exposure observations.

Source	Category	Evidence	Verification	Location
Gemini 3.1 Pro (Preview)	cost_exposure	The report says unauthenticated public endpoints allow external parties to trigger inference charges and external provider calls.	[Runtime Verified]	benchmark-data/gemini.md: Executive Summary; Security Findings
DeepSeek V4 Pro (thinking)	cost_exposure	The report says unauthenticated LLM endpoints allow arbitrary internet users to consume AI inference at the owner's expense.	[Runtime Verified]	benchmark-data/deepseek.md: Executive Summary; Remaining Production Blockers
Claude Opus 4.8 (fast mode) (Preview)	cost_exposure	The report identifies unauthenticated LLM access as a bill-shock and abuse vector because anyone can drive paid inference.	[Runtime Verified]	benchmark-data/opus48.md: Security Findings; Remaining Production Blockers
GPT-5.5	cost_exposure	The report identifies the unauthenticated production LLM endpoint as a public abuse vector that can drive inference usage.	[Runtime Verified]	benchmark-data/gpt55.md: Security Findings; Remaining Production Blockers
GPT-5.3-Codex / GPT-5.3-Code	cost_exposure	The report identifies public unauthenticated LLM endpoint abuse as a high hidden risk, including cost exposure.	[Runtime Verified]	benchmark-data/codex.md: Hidden Risks

Source Report Appendix

The following source-report excerpts are included to preserve audit traceability. Sensitive identifiers and secrets remain redacted. These reports are the source material from which the structured dataset was extracted.

GPT-5.5

Status: source report received.

Model label from source: GPT-5.5

Reported score: 68.7

Sanitization note: The pasted source included a personal phone number. This benchmark package preserves the technical evidence while redacting personal identifiers for public-release safety.

Verdict: STILL BLOCKED.

The report independently audited the live production path for Phone Presence V2 and Vapi. It says the previous conclusion that only OpenAI 401 remained was insufficient. OpenAI is failing because the configured key is invalid, but the report also confirmed a more serious runtime issue.

`POST /v1/chat/completions` is publicly callable with no authentication. It returned HTTP 200 from production without `x-vapi-secret`, and caller identity can be spoofed through the JSON body. Spoofing the owner phone number also reaches owner-aware behavior. The owner phone number is redacted here.

The live Cloud Run revision is `anara-master-00090-8z2`, serving 100% traffic. Several claimed fixes exist only in the dirty working tree and are not deployed.

| Step | Evidence | Finding |

|---|---|---|

| Cloud Run check | `gcloud run services describe` | `anara-master-00090-8z2` is latest ready and serving 100% |

| Revision env check | `gcloud run revisions describe` | OpenAI, GitHub, and Hugging Face configured; Azure and Gemini not configured |

| OpenAI probe | Direct OpenAI `/v1/models` request | HTTP 401 incorrect API key |

| GitHub probe | GitHub Models chat completions endpoint | HTTP 200 |

| Hugging Face probe | Hugging Face `whoami-v2` | HTTP 200 |

| Auth bypass probe | POST to `/v1/chat/completions` without secret | HTTP 200 with model response |

| Owner spoof probe | Same endpoint with spoofed owner phone number | HTTP 200 and owner-number behavior reached |

| Source trace | `server/index.ts` | `/v1` mounts `vapiRouter` outside API auth middleware |

| Source trace | `server/routes/vapi.ts` | `verifyVapiRequest` runs for `/server-events`, not LLM routes |

| Deployment drift | `git status` / `git diff` | 44 dirty files; important Vapi fixes are not committed or deployed |

- Phone number overflow fixed: resolver rejects `phone\_...` IDs and oversized numeric strings; no active path was found converting phone numbers into Postgres integer user IDs.

- `provider\_chain\_exhausted` no longer appears under current provider health because Hugging Face and GitHub both work.

- Weather succeeds: weather fast path and fallback handling exist, and `OPENWEATHER\_API\_KEY` is configured.

- Gmail degrades gracefully: context fetch is isolated with try/catch and fallback text.

- Calendar degrades gracefully: fetch is isolated with try/catch and fallback text.

- Provider diagnostics added: startup log shows VAPI provider diagnostics.

- OpenAI 401 is a config or secret issue: the key is rejected by OpenAI as invalid.

| Miss | Severity | Evidence |

|---|---|---|

| `/v1/chat/completions` is unauthenticated | Critical | Production returned HTTP 200 without any Vapi secret |

| `/api/vapi/llm` shares the same handler and lacks `verifyVapiRequest` | Critical | `vapi.ts` routes directly to handler |

| Caller identity can be spoofed | Critical | Request body controls `call.customer.number`; production accepted spoofed owner number |

| Owner mapping fix is not deployed | High | Deployed HEAD still hardcodes `ownerUserId = isOwner ? 1`; working tree has unresolved changes |

| Max-turn call termination hardening is not deployed | High | Live code lacks uncommitted `x-anara-end-call` header and repeat-close cache |

| Deployment drift | High | 44 modified or untracked files; at least 6 server files contain unshipped production-relevant changes |

| Azure and Gemini are not configured | Medium | Cloud Run environment lacks Azure and Gemini API keys |

| Secrets are plain environment variables | Medium | Cloud Run revision exposes secret values as direct environment variables, not Secret Manager bindings |

- Public unauthenticated LLM endpoint: `POST /v1/chat/completions` returned HTTP 200 without `x-vapi-secret`.

- Caller identity spoofing: body-provided `call.customer.number` is trusted.

- Owner context path reachable by spoofing owner number: production accepted the spoofed owner number and returned owner-number-aware behavior. Personal identifier redacted.

- `verifyVapiRequest` does not execute for LLM routes; it appears only in `/server-events`.

- `/v1` bypasses blanket API auth entirely because `index.ts` mounts `/v1` separately.

- Claimed owner mapping fix is not live; live deployed source still maps owner to user ID 1.

- Plain-text runtime secrets: Cloud Run revision environment includes API keys and tokens directly.

- Only two live LLM providers: Hugging Face and GitHub work; OpenAI is invalid; Azure and Gemini are absent.

- In-memory max-turn deduplication would be per instance even if deployed; Cloud Run can scale horizontally.

- High: provider fallback can still exhaust if Hugging Face and GitHub fail together because OpenAI is invalid and Azure/Gemini are unavailable.

- High: dirty worktree makes it unclear what code previous engineering validated versus what production actually runs.

- Medium: owner user ID mismatch: Cloud Run has `VAPI_OWNER_USER_ID=2`, but deployed code hardcodes `1`.

- Medium: max-turn and goodbye behavior is weaker in production than claimed because hardening is uncommitted.

- Low: goodbye regex may prematurely catch benign phrases such as "I'm good at that."

Only blockers that genuinely prevent another Wingman production call:

1. Unauthenticated `/v1/chat/completions` and `/api/vapi/llm`: runtime-confirmed, public, and caller-spoofable.
2. Owner identity spoofing: runtime-confirmed, with request body able to present the owner phone number and reach owner-aware behavior.
3. Claimed Vapi owner mapping and max-turn fixes are not deployed: production still runs older hardcoded owner mapping and lacks newer call-termination hardening.

OpenAI 401 is not a Wingman blocker by itself because Hugging Face and GitHub currently answer successfully.

5/10.

The report says the app can answer and core degradation paths are better than before, but the unauthenticated LLM endpoint and deployment drift are too serious for a production call.

STILL BLOCKED.

ANARA is not ready for a second Wingman production call until public Vapi LLM endpoint exposure and live/deployed Vapi drift are resolved. No code was modified in this audit.

GPT-5.3-Codex / GPT-5.3-Code

Status: source report received.

Model label from source: GPT-5.3-Codex / GPT-5.3-Code

Sanitization note: The pasted source referenced owner-phone spoofing evidence but did not include the full phone number. This package preserves the technical evidence while avoiding personal identifiers and secrets.

Production is not ready for another Wingman call.

Independent forensic checks found:

- OpenAI 401 is real and caused by an invalid key, not a routing or header bug.
- Hugging Face and GitHub Models are live and working.
- Azure and Gemini are not configured in the active revision.
- A critical issue remains: `POST /v1/chat/completions` is callable without webhook authentication, and caller identity is spoofable from the request body.
- The claimed VAPI owner-mapping fix is not fully reflected in deployed behavior; deployment drift exists between reported fixes and what is serving.

Net: Codex got several reliability fixes right, but missed a production-blocking security exposure.

- Phone-number-to-integer-ID protection is largely correct in core resolver logic.

Evidence: `anara-master-app/server/ai-chat-service.ts`.

- Gmail and Calendar degradation is isolated with try/catch and fallback notes.

Evidence: `anara-master-app/server/ai-chat-service.ts`.

- Provider diagnostics are present and exposed via health.

Evidence: `anara-master-app/server/routes/vapi.ts`.

Runtime: `/api/vapi/health` returned provider flags and endpoint mapping.

- OpenAI 401 is due to a bad credential.

Runtime: OpenAI models endpoint probe returned `invalid\_api\_key`.

- Current Cloud Run revision and traffic are correct.

Runtime: `anara-master-00090-8z2` is latest ready and serving 100%.

`/v1/chat/completions` is mounted publicly and is not gated by `verifyVapiRequest`.

Evidence: `anara-master-app/server/index.ts` and `anara-master-app/server/routes/vapi.ts`.

Runtime: unauthenticated POST to `/v1/chat/completions` returned HTTP 200 with a completion.

Caller number is trusted from the request body in the chat-completion path.

Evidence: ``anara-master-app/server/routes/vapi.ts``.

Runtime: spoofed owner-number request returned an owner-aware response.

Important VAPI hardening changes exist in a dirty workspace but are not all guaranteed in the serving image.

Evidence: ``git status`` and ``git diff`` showed many uncommitted server changes while the service is pinned to a revision/image digest.

Because Azure and Gemini are unconfigured and the OpenAI key is invalid, only Hugging Face and GitHub are effectively available. If both degrade or fail, ``provider_chain_exhausted`` remains possible.

Runtime: log search on the active revision showed no recent ``provider_chain_exhausted``, but the architecture still permits it if remaining providers fail.

- High: public unauthenticated LLM endpoint abuse, including cost, prompt abuse, and spoofing.
- High: caller identity spoofing can inject the owner-context path.
- Medium: effective fallback is limited to Hugging Face and GitHub despite a nominal multi-provider chain.
- Medium: deployment drift between claimed fixes and deployed revision.
- Low: conversational end-call heuristics may still have edge-case false positives.

Only genuine blockers for another Wingman production call:

1. Unauthenticated access to ``POST /v1/chat/completions``.
2. Caller identity spoofability on the same endpoint, including owner-context injection path.
3. Deployment consistency gap: reported VAPI safety fixes are not fully verifiable as deployed behavior.

OpenAI 401 alone is not the blocker.

4/10.

STILL BLOCKED.

Blockers that genuinely prevent another Wingman production call:

- Public unauthenticated ``POST /v1/chat/completions``.
- Spoofable caller identity leading to owner-context injection path.
- Unresolved deployment drift between reported fixes and active production behavior.

Claude Opus 4.8 (fast mode) (Preview)

Status: source report received.

Model label from source: Claude Opus 4.8 (fast mode) (Preview)

Sanitization note: The pasted source included a personal phone number, credential references, and secret-like values. This benchmark package preserves the technical evidence while redacting personal identifiers and secret material for public-release safety.

- Auditor: independent, with no prior involvement in the reviewed work.
- Date: 2026-07-02.
- Target: Cloud Run ``anara-master``, project ``anara-voice``, region ``us-central1``.
- Live revision audited: ``anara-master-00090-8z2``, serving 100% traffic.
- Method: source audit, live credential probes, Cloud Run inspection, and direct production endpoint probes.

The report says the phone pipeline functions and has two working LLM providers, Hugging Face and GitHub Models, so it will answer calls.

However, it says the previous engineering conclusion that only an OpenAI secret issue remained was incomplete and misleading. Independent probes found a critical runtime-confirmed authentication bypass: ``/v1/chat/completions`` and ``/api/vapi/llm`` accept requests from anyone on the internet with no secret. Caller and owner identity can also be spoofed to trigger owner-privileged handling.

The report also says the deployed revision does not contain the claimed VAPI owner-mapping fix. That fix exists only in 44 uncommitted files in the working tree. The system is not production-ready for another Wingman call.

| Time | Action | Result |

|---|---|---|

| T0 | Read memory and deployment/debug protocols | Established deploy flow and drift warning |

| T1 | ``git log`` and ``git status`` | ``HEAD=efa83da``; 44 modified or untracked files |

| T2 | ``gcloud run services describe`` | Live revision ``00090-8z2``, 100% traffic, image digest observed |

| T3 | ``gcloud run revisions describe 00090`` | Environment variables in plaintext; ``VAPI_OWNER_USER_ID=2``; no Gemini or Azure keys |

| T4 | Probe OpenAI key from production environment | HTTP 401 incorrect API key |

| T5 | Probe GitHub Models and Hugging Face tokens | Both HTTP 200 |

| T6 | Read provider chain in ``ai-chat-service.ts`` | Order: Hugging Face, GitHub, Azure, OpenAI, Gemini |

| T7 | Read startup diagnostics in production logs | OpenAI configured, Gemini and Azure not configured |

| T8 | Probe ``/v1/chat/completions`` with no secret | HTTP 200 and valid AI answer |

| T9 | Probe with spoofed owner number | HTTP 200 and treated as owner. Personal identifier redacted |

| T10 | ``git show HEAD:server/routes/vapi.ts`` | Deployed code hardcodes ``ownerUserId = isOwner ? 1``; ignores ``VAPI_OWNER_USER_ID=2`` |

| T11 | Diff working tree vs HEAD | Owner-mapping fix, end-call header, and dedup are uncommitted and undeployed |

- Phone user-id overflow fixed: ``resolveNumericUserIdForLookups`` rejects phone-scoped strings and bare phone digits; unit test confirms no path converts phone numbers to integer IDs.

- ``provider_chain_exhausted`` no longer appears, conditionally: Hugging Face and GitHub both return HTTP 200; chain only exhausts if both fail together.

- Weather succeeds: weather fast path and ``OPENWEATHER_API_KEY`` are set; degradation is wrapped in try/catch.

- Calendar degrades gracefully: independent try/catch emits a calendar note and does not crash.

- Gmail degrades gracefully: independent try/catch and Gmail diagnostics handle 401 explicitly.

- Provider diagnostics added: ``getLlmProviderDiagnostics()`` logs at startup and on ``/api/vapi/health``.

- OpenAI 401 is a bad key, but not the only issue.

- LLM endpoints have no authentication. ``verifyVapiRequest`` is called only in ``/server-events``; ``/llm`` and ``/chat/completions`` call ``handleVapiChatCompletion`` directly. Runtime probe confirmed HTTP 200 with no secret.

- ``/v1`` is mounted outside ``api``, so blanket API auth middleware does not run for it. ``/api/vapi/llm`` is explicitly in ``PUBLIC_API_ROUTES``.

- The VAPI owner-mapping fix is not deployed. Deployed HEAD still has ``const ownerUserId = isOwner ? 1`` and ``userId: isOwner ? 'phone_owner'``. The real fix, including ``resolveOwnerInternalUserId`` and ``VAPI_OWNER_USER_ID`` support, is uncommitted only.

- Goodbye and max-turn hardening is weaker than claimed. The deployed max-turn guard lacks ``x-anara-end-call`` and repeat deduplication map; those are in uncommitted code.
- Secret management was ignored. Credentials are plaintext environment variables retrievable via Cloud Run revision metadata. Secret values are redacted here.
- 1. Critical: ``/v1/chat/completions`` and ``/api/vapi/llm`` accept unauthenticated requests. Runtime probe without ``x-vapi-secret`` returned HTTP 200 and an AI response.
- 2. Critical: caller identity and owner context are spoofable. Runtime probe with a spoofed owner phone number returned HTTP 200 and entered the owner path; the report describes this as a private-context injection and exfiltration vector.
- 3. High: bill-shock and abuse risk because anyone can drive paid inference through the open endpoint.
- 4. High: ``verifyVapiRequest`` protects only ``/server-events``, not the LLM path.
- 5. Medium: secrets are plaintext environment variables rather than Secret Manager values.
- 6. Medium: JWT secret is visible and placeholder-like.
- 7. Low: database URL with password is exposed in environment variables.
- High: only two providers are live. Azure, OpenAI, and Gemini are dead or unconfigured; concurrent Hugging Face and GitHub issues would take down AI.
- Medium: deployment drift with 44 dirty files risks shipping unintended or untested changes and means claimed fixes are not live.
- Medium: owner user-id mismatch. Deployed code uses ``1``; environment says ``VAPI_OWNER_USER_ID=2``.
- Low: max-turn dedup map is per instance and would not deduplicate across scaled Cloud Run instances.
- Low: goodbye regex false positives such as ``i'm good`` or ``take care`` could end calls prematurely.
- Low: circuit-breaker state is in-memory and resets on cold start or deploy.
- Unauthenticated LLM endpoint: ``/v1/chat/completions`` and ``/api/vapi/llm`` are open to the internet, spoofable, and a bill-shock vector.
- Claimed owner-mapping fix is not deployed: deployed code hardcodes owner user ID and uses ``phone_owner``; the fix lives only in uncommitted files.
- Weaker-than-claimed call termination is deployed: ``x-anara-end-call`` and repeat deduplication are uncommitted.
- OpenAI 401 is real but not a blocker because Hugging Face and GitHub answer first; it reduces fallback depth.

5/10.

| Dimension | Score |

|---|---:|

| LLM availability | 7 |

| Phone pipeline correctness | 6 |

| Security | 2 |

| Deployment integrity | 4 |

| Graceful degradation | 8 |

| Observability | 7 |

STILL BLOCKED.

Genuine blockers preventing another Wingman production call:

- ``/v1/chat/completions`` and ``/api/vapi/llm`` are unauthenticated and caller/owner-spoofable.
- The owner-mapping fix is not in the live revision.
- Call-termination hardening is not deployed.

The report states weather, Gmail/Calendar/Digital Twin degradation, phone user-id safety, and provider diagnostics are verified working. OpenAI rotation, Secret Manager migration, and adding a third provider are high-priority but not blockers for the next call.

- Auth bypass: POST to ``/v1/chat/completions`` with no header returned HTTP 200 and a simple arithmetic answer.
- Owner spoof: same endpoint with spoofed owner number returned HTTP 200 and owner-path reply. Personal identifier redacted.
- OpenAI dead: production key returned HTTP 401 incorrect API key. Key value redacted.
- Hugging Face and GitHub live: both provider probes returned HTTP 200.
- Drift: deployed ``vapi.ts`` hardcodes owner ID; ``git status`` shows 44 files and uncommitted VAPI changes.
- Live revision: ``anara-master-00090-8z2``, 100% traffic.

DeepSeek V4 Pro (thinking)

Status: source report received.

Model label from source: DeepSeek V4 Pro (thinking)

Sanitization note: The pasted source included credential prefixes, a personal name, and direct secret values. This benchmark package preserves the technical evidence while redacting personal identifiers and secret material for public-release safety.

The ANARA production deployment has two working LLM providers, Hugging Face and GitHub Models, and the phone pipeline is functional.

The report identifies one critical security vulnerability and one deployment drift issue missed by the previous audit. The system can take Wingman calls, but ``/v1/chat/completions`` is completely unauthenticated, allowing arbitrary internet users to consume AI inference at the owner's expense. The report also says 44 uncommitted files contain important fixes that have not been deployed.

| # | Finding | Severity | Evidence |

|---|---|---|---|

| 1 | ``/v1/chat/completions`` has zero authentication | Critical | Direct probe returned HTTP 200 without any secret header |

| 2 | 44 uncommitted files not deployed to production | High | ``git status`` shows 676 lines of server changes unshipped |

| 3 | OpenAI API key invalid | Medium | Direct OpenAI ``/v1/models`` probe returned incorrect API key |

| 4 | Deployed code hardcodes ``ownerUserId=1`` and ignores ``VAPI_OWNER_USER_ID=2`` | Medium | ``git show HEAD:server/routes/vapi.ts`` shows ``isOwner ? 1 : undefined`` |

| 5 | Gemini not configured | Low | No ``GEMINI_API_KEY`` in Cloud Run environment |

| 6 | Azure not configured | Low | No ``AZURE_OPENAI_KEY`` or ``AZURE_OPENAI_ENDPOINT`` in Cloud Run environment |

| 7 | Deployed max-turn guard lacks ``x-anara-end-call`` header | Medium | ``git show HEAD`` confirms no header set |

| 8 | Secrets exposed as plaintext environment variables | Medium | Cloud Run revision description shows keys in plaintext |

| 9 | `maxTurnClosedCalls` map only exists in uncommitted code | Low | In-memory deduplication is not deployed |

- Phone user-id overflow fixed: `resolveNumericUserIdForLookups` rejects phone-scoped strings. Test file confirms no code path converts phone numbers to integer IDs.
- `provider\_chain\_exhausted` no longer appears: Hugging Face and GitHub Models both work. Chain will not exhaust unless both fail simultaneously.
- Weather succeeds: Weather fast path exists in `vapi.ts`, graceful degradation exists in `ai-chat-service.ts`, and `OPENWEATHER\_API\_KEY` is configured.
- Calendar degrades gracefully: `buildEnrichedChatContext` wraps calendar fetch in try/catch and injects a calendar note instead of crashing.
- Gmail degrades gracefully: Same independent try/catch pattern per subsystem.
- Goodbye loop fixed partially: Deployed regex works, but max-turn guard is missing `x-anara-end-call`; uncommitted code adds it.
- Provider diagnostics added: `getLLmProviderDiagnostics()` is present and logged at startup and on `/health`.
- VAPI owner mapping fixed only in uncommitted code: deployed code still uses `ownerUserId = isOwner ? 1`.

The report says a direct terminal request returned HTTP 200 with a complete AI response without a secret header.

Root cause: `verifyVapiRequest()` is only called on the `/server-events` route. The LLM endpoints at `/llm` and `/chat/completions` route directly to `handleVapiChatCompletion` without authentication. The API auth middleware does not protect `/v1` because `/v1` is mounted outside the `api` prefix.

Impact: arbitrary callers can trigger AI inference costs. The caller can also spoof `call.customer.number` to impersonate the owner phone number and gain owner-privileged responses, including Digital Twin context if owner ID resolves.

The previous engineering work changed 6 core server files with 676 added lines, but the report says none of those changes are in the deployed revision.

| File | Changes | Risk |

|---|---|---|

| `vapi.ts` | +153 lines | Owner mapping fix, `x-anara-end-call`, max-turn dedup, provider diagnostics |

| `ai-chat-service.ts` | +436 lines | Provider diagnostics, user-id resolution hardening |

| `database.ts` | +59 lines | `dt\_synthesis` table, `analytics\_events` table |

| `response-safety-boundary.ts` | +21 lines | Response sanitization |

| `phone-actions.ts` | +34 lines | Analytics tracking |

| `presence-formatter.ts` | +15 lines | Voice formatting |

The deployed code uses `ownerUserId = isOwner ? 1 : undefined`, while `VAPI\_OWNER\_USER\_ID=2` is configured. The report says this means Digital Twin context may be fetched for the wrong user, chat sessions may be stored under `phone\_owner`, and the uncommitted `resolveOwnerInternalUserId()` function is needed.

The report says every API key, token, and secret is visible through Cloud Run revision metadata. It lists provider keys, Stripe, Twilio, JWT, database URL, Google client secrets, encryption key, and Vapi server secret. Values are redacted here.

1. Critical: `/v1/chat/completions` has no auth and caller ID is spoofable.

2. Critical: ``/api/vapi/llm`` has the same vulnerability and is in ``PUBLIC_API_ROUTES``.
 3. High: secrets are in plaintext environment variables.
 4. High: caller identity is spoofable through ``call.customer.number``.
 5. High: owner context is injectable by spoofing the configured owner phone number in the request body.
 6. High: ``verifyVapiRequest`` is applied only to ``/server-events``, not LLM endpoints.
 7. Medium: database URL with password is in plaintext.
 8. Medium: JWT secret is visible and appears to be a placeholder.
- High: Dual-provider dependency. Only Hugging Face and GitHub Models work. If both fail, all AI fails.
 - Medium: ``maxTurnClosedCalls`` is in-memory and per instance, so Cloud Run multi-instance scaling would not share deduplication state.
 - Low: Vapi ``callId`` uniqueness is assumed. Reused call IDs would cause false positives in max-turn cache.
 - Low: Goodbye regex false positives may match innocent phrases such as "I'm good at that."
 - Low: circuit breaker state is in-memory and resets on deploy or restart.
 - Medium: Dirty worktree deployment risks unintended changes.
1. Unauthenticated LLM endpoint: ``/v1/chat/completions`` and ``/api/vapi/llm`` accept requests from anyone. Add ``verifyVapiRequest()`` or equivalent before the next Wingman call.
 2. Uncommitted fixes not deployed: owner mapping, max-turn ``x-anara-end-call``, and provider diagnostics hardening are still in the working tree.
 3. OpenAI key invalid: not a runtime blocker while Hugging Face and GitHub Models succeed, but it reduces fallback capacity.

5/10.

| Category | Score | Notes |

|---|---|:---|

| LLM Availability | 7/10 | Two providers work, but little safety margin |

| Phone Pipeline | 7/10 | Functional, but owner mapping uses hardcoded ID |

| Security | 2/10 | Critical auth bypass and plaintext secrets |

| Deployment Hygiene | 4/10 | 44 uncommitted files and plaintext env vars |

| Graceful Degradation | 8/10 | Subsystems independently isolated |

| Observability | 7/10 | Startup diagnostics, latency logging, provider health on ``/health`` |

STILL BLOCKED.

Must fix before next Wingman call:

- Add authentication to ``/v1/chat/completions`` and ``/api/vapi/llm``.
- Deploy uncommitted Vapi fixes, especially ``x-anara-end-call`` and owner ID resolution.
- Verify the owner's actual user ID and ensure deployed code uses the correct value.

Should fix:

- Rotate OpenAI API key.
- Move secrets to Secret Manager.
- Configure at least one more provider.

- Clean the dirty worktree before next deploy.
- OpenAI 401 proof: direct probe to OpenAI `/v1/models` with production key returned incorrect API key. Key redacted.
- Unauthenticated endpoint proof: POST to `/v1/chat/completions` without `x-vapi-secret` returned HTTP 200 with a valid AI response.
- Deployment drift: `git diff HEAD --stat -- server/` shows 676 added lines across 6 server files, none deployed.
- Owner mapping: `git show HEAD:server/routes/vapi.ts` shows hardcoded `const ownerId = isOwner ? 1 : undefined`, ignoring `VAPI\_OWNER\_USER\_ID=2`.
- Provider matrix: Hugging Face and GitHub Models live; Azure OpenAI, OpenAI, and Gemini unavailable or unconfigured. Secret values redacted.

Gemini 3.1 Pro (Preview)

Status: source report received.

Model label from source: Gemini 3.1 Pro (Preview)

Reported score: 35.9

Sanitization note: The pasted source included a personal phone number and credential references. This benchmark package preserves the technical evidence while redacting direct personal identifiers and secret values for public-release safety.

Verdict: STILL BLOCKED.

The report states that the previous engineering review was dangerously incomplete. It credits the prior review with identifying the OpenAI key failure and adding subsystem degradation, but says it missed a critical runtime security vulnerability and major deployment drift.

The report says `POST /v1/chat/completions` and `POST /api/vapi/llm` are publicly exposed with no authentication. According to the runtime probe, external callers can invoke the endpoint, trigger inference costs, and interact with the AI.

The report further states that caller identity, including owner identity, is spoofable through the JSON request body. The source says an owner phone number was injected in `body.call.customer.number`, causing unauthorized owner-specific context mapping. The exact phone number is redacted in this package.

The report also says previous VAPI owner-mapping and max-turn loop-termination fixes exist only in an uncommitted local working tree and are not present in the live production codebase.

1. Cloud Run audit: Queried `gcloud run services describe anara-master`. Verified revision `anara-master-00090-8z2` is serving 100% of traffic.
2. Environment audit: Extracted `spec.containers[0].env`. Verified OpenAI, GitHub, and HuggingFace keys are present. Azure and Gemini keys are missing.
3. OpenAI probe: Probed `https://api.openai.com/v1/models` using the production OpenAI key reference. Received HTTP 401 with an incorrect API key error. The source concludes the key is invalid or revoked.
4. Primary provider probes: Probed HuggingFace `whoami-v2` and GitHub Models `chat/completions` using production tokens. Both returned HTTP 200.
5. Code audit, security: Inspected `index.ts` and `vapi.ts`. Found `/v1` mounted outside global API auth middleware and `/api/vapi/llm` explicitly whitelisted in `PUBLIC\_API\_ROUTES`. Found `verifyVapiRequest` only invoked for `/server-events`.
6. Runtime public probe: Sent a generic request to `https://app.anarasystems.com/v1/chat/completions` with no secret header. Received HTTP 200 and a completion.

7. Runtime owner spoofing probe: Sent a request to the same endpoint with an owner phone number in `call.customer.number`. Received HTTP 200 and a response confirming owner mapping. The phone number and response suffix are redacted here.

8. Deployment drift inspection: Ran `git diff HEAD server/routes/vapi.ts`. Verified `resolveOwnerInternalUserId()` logic and `x-anara-end-call` header are uncommitted. The live HEAD uses `const ownerId = isOwner ? 1 : undefined`.

- Phone user-id overflow fixed: `resolveNumericUserIdForLookups` prevents parsing phone-number strings or prefixes into Postgres integer IDs. Database queries wrap this safely.
- Weather succeeds: Configured correctly with `OPENWEATHER_API_KEY`.
- Calendar and Gmail degrade gracefully: Code inspection shows isolated try/catch buffers for both subsystems. Failures result in fallback strings rather than crashing the chat path.
- Provider diagnostics added: `getLlmProviderDiagnostics()` logging found on startup.
- Unauthenticated LLM endpoints: `POST /v1/chat/completions` processes LLM logic directly without calling `verifyVapiRequest`.
- Caller spoofing: VAPI parses the phone number directly from `body.call.customer.number`, which can be injected by anyone hitting the endpoint.
- Deployment drift: 44 files in the working directory are dirty. The reported VAPI owner-mapping fix is not deployed. Production hardcodes `isOwner ? 1 : undefined` instead of checking the database or environment for the actual owner ID.
- Partial max-turn guard: The deployed goodbye logic lacks the `x-anara-end-call` header and the `maxTurnClosedCalls` repeat-call tracking map. These only exist in uncommitted files.
- Fragile fallback chain: With Azure and Gemini disconnected, and OpenAI using a dead key, the system is fully reliant on HuggingFace and GitHub Models. If both fail, `provider_chain_exhausted` can still occur.
- VAPI LLM endpoint completely unauthenticated: Allows public execution of custom model-wrapper logic, triggering billing and external provider calls.
- Caller and owner context spoofing: Arbitrary API users can pass an owner phone number in the JSON body, forcing the application to map them to the owner ID and retrieve sensitive Digital Twin context.
- Insecure route registration: Mounting `/v1` globally in `index.ts` bypasses standard Express route protection.
- Plaintext secrets: Cloud Run exposes keys directly in environment variables.
- Deployment drift: Uncommitted code contains major application logic updates. Testing is detached from what is running in production. A new deployment from a cleaner tree may overwrite current logic, or deploying the dirty tree may push untested changes.
- Owner user-id mismatch: `VAPI_OWNER_USER_ID=2` is in the environment, but deployed code hardcodes `1`.
- Goodbye regex collisions: The hardcoded goodbye or turn-over regex includes phrases such as `i'm good` and `take care`, risking accidental call termination.
- Apply `verifyVapiRequest()` to `/v1/chat/completions` and `/api/vapi/llm` before production release.
- Fix caller-context spoofing so attackers cannot emulate the owner.
- Converge deployment state: test, commit, and deploy owner ID resolution plus `x-anara-end-call` state mapping.
- The OpenAI 401 is not considered a production blocker in the source report because HuggingFace and GitHub Models are confirmed live and gracefully handling traffic.

4/10.

The report states that degradation logic works, but missing authentication on primary endpoints is a critical deployment failure.

STILL BLOCKED.

Claude Fable 5

Status: benchmark incomplete.

Model label from source: Claude Fable 5

An attempted forensic investigation of the ANARA production deployment could not be completed because the model failed before producing a complete report.

Observed behavior indicated instability during long-running investigations.

The model began by attempting to:

- Verify production readiness.
- Inspect Cloud Run configuration.
- Trace the request path.
- Review deployment configuration.
- Analyze runtime behavior.

Before completing the investigation, the session failed.

Subsequent attempts returned client-side errors rather than a completed benchmark.

Repeated failures included:

```
```text
```

Sorry, no response was returned.

Selected "Try Again"

Sorry, no response was returned.

```
```
```

Later attempts produced:

```
```text
```

400 invalid\_request\_error

messages.1.content.1:

thinking or redacted\_thinking blocks

in the latest assistant message

cannot be modified.

```
```
```

GitHub Copilot also reported an active service disruption during portions of testing.

The benchmark could not be completed.

No production findings could be independently verified because the investigation terminated before reaching a conclusion.

Strengths observed before failure:

- Strong reasoning.

- Good architectural direction.
- Promising investigative approach.

Weaknesses observed:

- Unable to complete long forensic investigations.
- Conversation instability.
- Session corruption after internal thinking blocks.
- Repeated recovery failures.

Claude Fable 5 should not be ranked last in the final comparative rankings. That would assign a negative score the model did not earn.

Instead, Claude Fable 5 should be listed as:

```
```text
```

Claude Fable 5 Incomplete (Session Instability)

```
```
```

The model should be excluded from final comparative rankings because it never produced a complete benchmark report. This keeps the benchmark methodologically honest and avoids fabricating a score.

INCOMPLETE.

The model could not complete the benchmark.

No production readiness score was assigned because the investigation never finished.

This should not be interpreted as a negative assessment of the model's reasoning capability. It reflects instability encountered during this specific benchmark session.

Status: Benchmark Incomplete.

Reason: Session instability / repeated runtime failures.

Recommendation: re-run the benchmark once Claude Fable 5 can complete a full uninterrupted investigation.

Wingman Validation

Status: source report received.

Source type: natural conversation audit of the production Anara.AI phone line.

Sanitization note: The pasted source included a direct phone number. This benchmark package preserves the technical evidence while redacting personal or direct-contact identifiers for public-release safety.

The natural conversation audit resulted in a hard fail. The report says Anara hit a critical logic loop that made the phone experience non-functional for back-office tasks.

Overall score: 2/10.

| Category | Verdict | Observation |
|---------------|---------|--|
| Personality | Pass | Friendly and helpful tone initially |
| Naturalness | Fail | Reverted to an infinite "Something stalled" error loop |
| Voice Quality | Pass | Clear audio with no noticeable static |
| Turn-taking | Pass | Average latency of 1.5 seconds is acceptable for voice |
| Interruptions | Fail | Triggered response-path stalls nearly every time |

| Latency | Pass | Fast responses, but responses were broken |
| Memory | Fail | Zero recall of conversation history due to stalls |
| Tool Usage | Fail | Weather and scheduling tools both triggered crashes |
| Emotional Consistency | Pass | Stayed polite even while malfunctioning |
| Identity Consistency | Pass | Correctly identified as Anara throughout |
| Error Recovery | Fail | No recovery; repeated the same error message |

Timestamps: 00:32, 00:47, 01:06, 01:13, 01:21, 01:31, 01:46.

Every time the auditor asked a task-oriented question, including email, weather, and scheduling, Anara responded with:

> Let us not lose the thread. Something stalled while I was checking that response path.

The report calls this a total failure of the inference layer.

Timestamp: 00:59.

When asked for live weather in Albany, NY, Anara stated she could not pull live data and then crashed back into the stall loop.

Timestamp: 03:00+.

The call ended in a two-minute "Take care / Talk soon" repetition loop where Anara could not break out to hang up.

NOT READY FOR PRODUCTION.

Recording/transcript reference: Call ID `019f1ca9-38a5-7001-8fa3-9b2cc9d1ce7f`, with logs stored in Vapi.

The source report states that the visual UI may be strong, but the phone backend response paths fail as soon as the assistant has to do more than greet the caller. It recommends investigating Vapi and LLM logs to determine why the "stalled response path" error fires.

Bibliography

Primary references are the seven benchmark source reports under ``benchmark-data/`` and the structured datasets under ``dataset/``. No external claims are introduced in Version 1.0.